


GameHorizon Suite: Multi-Horizon Data and Evaluation in Gameplay

Yiran Wang, Xingyilang Yin, Junfu Pu, Guangzhi Wang, Kaifeng Li, Mingyu Ouyang, Huiqiang Sun, Lingen Li, Cheng Cheng, Wangbo Yu, Honghao Chen, Xiaodong Cun, Chi-Man Pun, Zhiguo Cao, Ying Shan

<https://arxiv.org/abs/2609.25001>



Abstract

Modern video games provide a measurable testbed for AI models, combining abilities of visual understanding, instruction decomposition, goal planning, and precise action control over multiple temporal horizons. Existing datasets and benchmarks, however, either cover a narrow range of games, lack language instructions, or rely on high-variance online rollouts. To address these challenges, we introduce GameHorizon, a unified data and evaluation suite that measures gameplay capabilities at different horizons for diverse model families. GameHorizon Suite consists of three components. First, GameHorizon-Annotator is a scalable and automated annotation pipeline for multi-horizon instructions. Second, utilizing the pipeline, we construct GameHorizon-Data, the first large-scale AAA gameplay dataset with temporally aligned videos, player actions, and multi-horizon instructions. It comprises 5,000 hours of recordings from 21 games, collected by 100 human expert players. Third, we build GameHorizon-Bench with reproducible offline and stepwise online testing. The offline track enables reproducible evaluation using thousands of standardized questions organized into three primary tasks and a series of diagnostic variants, while the online track tests whether offline scores reflect actual gameplay capabilities and localizes failures to specific steps within long-horizon gameplay. Based on our GameHorizon Suite, we evaluate 47 models through more than one million model invocations, revealing a meaningful hierarchy of task difficulty and pronounced differences in model capabilities. Our work can provide a standardized yardstick for evaluating gameplay capabilities across horizons and model families. We will release our dataset, annotator, and benchmark to facilitate future research.

Task

Measuring Multi-Horizon Gameplay

- The paper introduces GameHorizon, a unified data and evaluation framework.
- The corpus comprises 5,000 hours of human gameplay across 21 game titles.
- The suite evaluates long-horizon gameplay through order-dependent causal tasks and order-flexible thematic tasks.
- Each task contains 2-6 verifiable subtasks to assess short-horizon capabilities.

Challenges

Narrow Coverage and Reproducibility Issues

- Existing datasets like GameWorld target only simple mini-games, resulting in narrow game coverage.
- Prior methods use small samples for low-confidence comparisons and online rollouts that are hard to reproduce.
- Earlier work focuses on Minecraft or non-AAA games.

Proposed method

Annotator Pipeline and Benchmark Design

- Action-aware segmentation uses keyboard-mouse traces to identify key action transitions and determine clip boundaries.
- Dynamic programming enforces duration ranges: 1–5 seconds for L1, 1–2 minutes for L2, and 5–8 minutes for L3.
- Unified multiple-choice questions ensure standardized evaluation where models directly output selections from four options.

Main results

Proprietary Models Lead, UMMs Lag

- Mean accuracy is 64.7%, substantially above the 25% random baseline.
- Proprietary model GPT-6-Astra ranks first with 80.2% overall accuracy.
- Best UMM BAGEL-7B-MoT achieves 61.4%, falling into Tier 3 or Tier 4.